MANO HARSHA DAVULURI

+91 9390003039 | manoharsha44@gmail.com | www.linkedin.com/in/mano-harsha-davuluri | Bengaluru, India

PROFESSIONAL SUMMARY

Results-driven **NITian** and **AI & Data Engineer** with 3+ years of experience building production Generative AI applications and scalable data platforms. Specializes in agentic AI (MCP, LLM tool-calling, RAG with vector search), LLMOps with MLflow, and FastAPI services on Databricks, with a strong foundation in Spark, AWS, data lakes, and ETL pipelines. Architected and launched an enterprise conversational AI assistant that enables natural-language data discovery across data assets, products, and domains. Proven track record of leading cross-functional teams and driving AI-led innovation with data governance at enterprise scale.

TECHNICAL SKILLS

Generative AI & Agents: Agentic AI, Multi-agent Systems (Agent-to-Agent), RAG, MCP (FastMCP), LangChain, LangGraph, ChatDatabricks, LLMs (Claude, Llama), Prompt Engineering, Guardrails, Embeddings, Hybrid Search (pgvector HNSW, RRF).
LLMOps & AI Platforms: MLflow 3 (Tracing, LLM Judges), Model Serving, AI Gateway (Routing, Fallbacks), Databricks Apps.
Cloud & Platforms: Databricks (Unity Catalog, Lakebase, Delta Lake, SQL Warehouses), AWS (S3, Lambda, EKS/Kubernetes).
Data Engineering: Apache Spark, PySpark, Spark SQL, Delta Lake, Brickflow, SparkExpectations, DBT, Informatica CDGC.
Data Modeling: Dimensional Modeling, Star/Snowflake Schema, SCD Type 1/2, Medallion, Lakehouse, Data Governance.
Programming & DevOps: Python, SQL, FastAPI Microservices, Git, Docker, Jenkins CI/CD, OAuth 2.0 (Okta, M2M), pytest.

PROFESSIONAL EXPERIENCE

Nike | *Software Engineer II – AI* — *Bengaluru, India | Jan 2026 – Present*

Compass AI – Agentic chatbot for Nike's enterprise data catalog

- Developed and extended an **MCP (Model Context Protocol)** server on Databricks Apps for a **LangChain AI agent** powered by **Claude Sonnet 4.5 (Anthropic)**, growing it to **12 LLM-callable tools**, incl. a **multi-agent** (agent-to-agent) tool calling Nike's AirTalk chatbot on the user's behalf via **Okta OAuth**, and tools for AI-proposed data products, docs lookup, and catalog stats.
- Created retrieval tools over **Lakebase (Postgres)** with **hybrid semantic search** (pgvector + full-text via **Reciprocal Rank Fusion**) and safe SQL access; replaced an external REST API with direct table queries, removing a network hop and its auth.
- Instrumented end-to-end **MLflow 3 tracing**, propagating request/session IDs from FastAPI to the MCP server via HTTP headers to group each multi-turn chat into one session, scored by **LLM-as-a-judge** evaluators (completeness, role adherence, tool-call efficiency, user frustration) with time-to-first-token monitoring.
- Configured **Databricks AI Gateway** with weighted **traffic splitting** (70% Claude Sonnet 4.6, 30% Claude Opus 5.5) and session-sticky routing, plus an automatic **fallback** to Llama 3.3 70B for high availability; added Anthropic **prompt caching** to cut token cost and latency.
- Crafted the agent's **system prompt and guardrails**: scope restriction to Nike metadata, non-disclosure of internals, correct bot identity, and accurate deep links to products, assets, and metrics.
- Productionized the MCP server across **dev/QA/prod** with **Jenkins CI/CD**, per-env config, **service-principal OAuth M2M**, and Cerberus secrets; retired a legacy MCP server and scheduled access-scope and vector index refresh jobs.
- Shipped **FastAPI** chatbot features: chat-history APIs (sessions, timestamped recent queries, per-session limits), a live health endpoint for UI "Online" status, and resilient ServiceNow ticketing, at a **90% test coverage** gate.
- Built medallion **ETL pipelines** (Spark, Brickflow) and a gold-layer **dimensional model** (22 dims, 5 facts) with SparkExpectations DQ rules, grounding the agent across **100,000+** data assets.

Epsilon (Publicis Groupe) | *Data Engineer* — *Bengaluru, India | Aug 2023 – Dec 2025*

- Designed a comprehensive Data Quality framework for the Prospect project, featuring automated discrepancy detection and real-time alerting mechanisms to ensure data integrity and compliance.
- Constructed scalable ETL/ELT data solutions processing over 150+ billion records daily into data lakes and warehouses using Spark, and trained ML models that enhanced decision-making, driving a 25% boost in quarterly revenue.
- Optimized and tuned Spark jobs for large-scale datasets (30+ TB), applying performance best practices and reducing infrastructure costs by 46%.
- Partnered with cross-functional teams to define data governance standards and set up CI/CD pipelines for data workflows.
- Owned end-to-end data pipelines spanning ingestion, extract-transform-load processes, data archiving, warehousing and recurring reporting using Apache Spark on Databricks Delta tables.
- Worked closely with solution architects and business stakeholders to design and deploy complex ML models for predictive analytics and BI solutions.

CERTIFICATIONS

- Databricks Certified Generative AI Engineer Associate – Databricks (Issued 2026)
- Databricks Certified Data Engineer Associate – Databricks (Issued 2024)

EDUCATION

National Institute of Technology (NIT) Karnataka | *Aug 2019 - May 2023*
Bachelor of Technology | Surathkal, Karnataka